

Function multiversioning in system image

All about vector registers

Yichao Yu

Harvard

Oct. 16, 2017

Implement function multi versioning in sysimg #21849



yuyichao wants to merge 9 commits into `master` from `yyc/codegen/clone`



Conversation 99



Commits 9



Files changed 30



yuyichao commented on 13 May • edited

Member



Current TODO:

- ☒ Implement more flexible cloning check
- ☒ Handle "PLT" cloning
- ☒ Generate dispatch data
- ☒ Update sysimg loading
- ☒ Fix `code_native`
- ☒ ? Figure out what's wrong with AVX on x86.

Goal

Amount of implementation details

- Just enough to make Jeff happy
- Without losing the audience
- With a single version of the talk (**Important!**)

Goal

Amount of implementation details

- Just enough to make Jeff happy
- Without losing the audience
- With a single version of the talk (**Important!**)

Goal

Amount of implementation details

- Just enough to make Jeff happy
- Without losing the audience
- With a single version of the talk (**Important!**)

Goal

Amount of implementation details

- Just enough to make Jeff happy
- Without losing the audience
- With a single version of the talk (**Important!**)

No performance scaling in threading #17395



ranjanan opened this issue on 13 Jul 2016 · 41 comments



ranjanan commented on 13 Jul 2016

Member




Consider the following code:

```
using Base.Threads
println("Number of threads = $(nthreads())")
x = rand(10^6)
y = zeros(10^6)
println("Warmup!")
for i = 1:10^6
    y[i] = sin(x[i])^2 + cos(x[i])^2
end
t1 = @elapsed for i = 1:10^6
    y[i] = sin(x[i])^2 + cos(x[i])^2
end
@assert sum(y) == 10^6
t2 = @elapsed @threads for i = 1:10^6
    y[i] = sin(x[i])^2 + cos(x[i])^2
end
@assert sum(y) == 10^6
println("Serial time = $t1")
println("Parallel time = $t2")
```

How did it all begin

No performance scaling in threading #17395

Open ranjanan opened this issue on 13 Jul 2016 · 41 comments



Register size

ax → **eax** → **rax**

mmx0 → **xmm0** → **ymm0** → **zmm0**

Register size

ax → **eax** → **rax**

mmx0 → **xmm0** → **ymm0** → **zmm0**

More/larger registers → More states to manage

Register size

ax → **eax** → **rax**

mmx0 → **xmm0** → **ymm0** → **zmm0**

More/larger registers → More states to manage

(Intel got something wrong almost everytime)

Compile openlibm for AVX (ymm registers)

- GCC function attribute *target_clone*
...specify that a function be cloned into multiple versions compiled with different target options ... It also creates a resolver function (...) that dynamically selects a clone suitable for current architecture ...
- Has bug on current GCC versions so not usable in openlibm.
But it's a good start.

Compile openlibm for AVX (ymm registers)

- GCC function attribute *target_clone*
...specify that a function be cloned into multiple versions compiled with different target options ... It also creates a resolver function (...) that dynamically selects a clone suitable for current architecture ...
- Has bug on current GCC versions so not usable in openlibm.
But it's a good start.

Why function multiversioning

- For code in the system image
- For JIT code

Why function multiversioning

- For code in the system image
- For JIT code

Why function multiversioning

- For code in the system image
 - For JIT code
- JIT features > system image features

Why function multiversioning

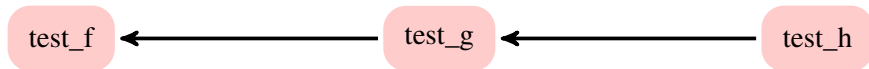
- For code in the system image

- For JIT code


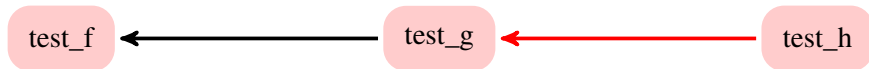
JIT features > system image features

JIT register size == system register size


Cloning



Cloning



Cloning



Next step

- Expose user API
 - ▶ Platform agnostic, e.g. “fma”
 - ▶ Platform specific, e.g. “x86::avx2”
- Make use of this in other part of runtime

Next step

- Expose user API
 - ▶ Platform agnostic, e.g. “fma”
 - ▶ Platform specific, e.g. “x86::avx2”
- Make use of this in other part of runtime

Next step

- Expose user API
 - ▶ Platform agnostic, e.g. “fma”
 - ▶ Platform specific, e.g. “x86::avx2”
- Make use of this in other part of runtime

Next step

- Expose user API
 - ▶ Platform agnostic, e.g. “fma”
 - ▶ Platform specific, e.g. “x86::avx2”
- Make use of this in other part of runtime

Next step

- Expose user API
 - ▶ Platform agnostic, e.g. “fma”
 - ▶ Platform specific, e.g. “x86::avx2”
- Make use of this in other part of runtime
crc32c, PLT, ...

